

Hybrid Adaptive DNA Indexing Engine with Learned Search Optimization

1st Sonali Antad

Dept. Computer Engineering

Vishwakarma Institute of Technology

Pune-411037, India

sonali.antad@vit.edu

2nd Lakshit Barekar

Dept. Computer Engineering

Vishwakarma Institute of Technology

Pune-411037, India

lakshit.barekar24@vit.edu

3rd Tejas Bankar

Dept. Computer Engineering

Vishwakarma Institute of Technology

Pune-411037, India

tejas.bankar24@vit.edu

4th Anushka Auti

Dept. Computer Engineering

Vishwakarma Institute of Technology

Pune-411037, India

anushka.auti24@vit.edu

5th Yash Bhandari

Dept. Computer Engineering

Vishwakarma Institute of Technology

Pune-411037, India

yash.bhandari24@vit.edu

Abstract—With the exponential growth of genomic data, efficient indexing and querying of DNA sequences have become critical challenges in bioinformatics. Traditional data structures such as suffix arrays and FM-indexes provide reliable solutions but suffer from scalability and performance limitations. This paper proposes a Hybrid Adaptive DNA Indexing Engine that integrates classical suffix array techniques with learned indexing models and intelligent error prediction mechanisms. By combining adaptive data structures with machine learning-based prediction, the proposed system significantly reduces query time while maintaining high accuracy. Experimental analysis demonstrates improved performance in large-scale genomic datasets, making the system suitable for real-time DNA sequence analysis and pattern matching applications.

Index Terms—DNA Indexing, Suffix Array, Learned Index Structures, FM-Index, Bioinformatics, Pattern Matching.

I. INTRODUCTION

The rapid advancement of high-throughput sequencing technologies, particularly next-generation sequencing (NGS), has led to an unprecedented growth in genomic data. Modern sequencing platforms are capable of generating terabytes of DNA sequence data in a single experiment, creating significant challenges in terms of storage, indexing, and efficient retrieval. As a result, the development of scalable and high-performance algorithms for genomic data processing has become a critical area of research in bioinformatics [6], [15].

DNA sequence analysis is fundamental to a wide range of biological and medical applications, including disease diagnosis, personalized medicine, genome assembly, mutation detection, and evolutionary studies. Many of these applications rely heavily on efficient pattern matching techniques, where a query sequence must be located within a large reference genome. This problem, commonly referred to as substring search in DNA sequences, becomes computationally intensive due to the massive size and repetitive nature of genomic data [7], [9].

Classical data structures such as suffix arrays and FM-indexes have been extensively used to address this problem. Suffix arrays provide a space-efficient alternative to suffix trees while supporting fast pattern matching operations, whereas FM-indexes leverage the Burrows–Wheeler Transform to achieve compressed indexing [12], [14]. Despite their effectiveness, these traditional approaches face several limitations. They often incur high memory overhead when scaling to large genomes, suffer from cache inefficiencies due to irregular memory access patterns, and may not fully utilize modern hardware architectures [2], [4].

To overcome these challenges, recent research has explored the integration of machine learning techniques with classical indexing structures. Learned index structures treat indexing as a prediction problem, where models approximate the position of data within a sorted structure. This paradigm has shown promising results in reducing search time and improving cache locality in various domains [13]. In genomic applications, approaches such as Sapling demonstrate how learned models can enhance suffix array performance by narrowing the search space and reducing query time [1].

In this context, the present work proposes a hybrid architecture that combines suffix array–based indexing with adaptive learning models and intelligent error prediction mechanisms. The learning component estimates the position range of query patterns within the suffix array, thereby significantly narrowing the search space. Additionally, error prediction techniques are employed to correct model inaccuracies, ensuring both efficiency and accuracy. By leveraging the strengths of both classical data structures and modern machine learning approaches, the proposed system aims to achieve improved performance in terms of query time, memory utilization, and scalability.

Figure 1.1 illustrates the flowchart of the proposed DNA sequence search system. The process begins with the input of genomic DNA sequences along with the query pattern to be searched. The input data is then passed through a preprocessing stage, where the sequence is cleaned, normalized, and encoded for efficient computation. Next, a suffix array is constructed from the processed DNA sequence, enabling structured and sorted indexing of all possible suffixes. The system then utilizes a learning module, where a predictive model is trained to estimate the approximate position of query patterns within the suffix array. Based on this prediction, the search module narrows down the search space and performs a bounded search instead of scanning the entire dataset. A verification step ensures that the identified matches are accurate through character-by-character comparison. Finally, the output module returns the matched positions along with performance metrics such as execution time and search efficiency. This flow demonstrates how the integration of classical indexing and predictive modeling significantly enhances the overall performance of DNA sequence search.




Fig 1.1 Flowchart of working project

II. LITERATURE REVIEW

Efficient DNA sequence analysis has become increasingly important due to the rapid growth of genomic data. One of the fundamental challenges is designing scalable indexing and search techniques for large biological datasets. Kirsche *et al.* [1]

introduced Sapling, a learned index-based approach that enhances suffix arrays using machine learning models. By predicting approximate suffix positions, the approach reduces binary search space and improves query speed significantly, especially for large genomes. Similarly, Wu [2] explored enhanced suffix arrays (ESA) with bitpacking and compression techniques, achieving improvements in both memory efficiency and search speed for genomic queries.

Barsky *et al.* [3] proposed the Suffix Rank algorithm, which supports scalable suffix array construction for datasets exceeding main memory capacity. Their work focuses on reducing disk I/O overhead using sequential access patterns. Further improvements were introduced by Rabea *et al.* [4], who optimized suffix array construction specifically for DNA alphabets (A, C, G, T), reducing computational complexity by leveraging domain-specific constraints. Similarly, Yang *et al.* [5] developed a k-ordered FM-index construction technique using induced sorting, which enhances both query performance and memory efficiency in genomic indexing.

In addition to suffix-based methods, alternative indexing strategies have been explored. Wohoush and Saheb [6] proposed wavelet transformation-based indexing for large DNA databases, which reduces I/O operations and improves search efficiency by transforming sequences into numerical representations. Their approach highlights the importance of selecting appropriate index structures and window sizes to balance space and search performance. Other techniques such as dictionary-based indexing, multi-resolution structures, and clustering-based indexing have also been studied, though they often suffer from high space complexity or false matches.

String matching algorithms form the backbone of DNA search operations. The classical Knuth–Morris–Pratt (KMP) algorithm [7] provides efficient linear-time exact pattern matching using prefix functions, making it suitable for large-scale DNA sequence searches. Additionally, Boyer–Moore and Rabin–Karp algorithms [8] have been widely used for efficient substring search by reducing unnecessary comparisons.

Approximate string matching is equally important in genomic analysis, especially for mutation detection. Levenshtein distance-based approaches [9] measure similarity between sequences by counting insertions, deletions, and substitutions. These methods are often combined with heuristic techniques such as BLAST to improve performance while maintaining accuracy in large datasets.

Biological characteristics of DNA sequences also influence computational methods. Benita *et al.* [10] demonstrated that regional GC content is a strong predictor of PCR success, emphasizing the importance of preprocessing DNA sequences for better analysis. Similarly, Blackburn *et al.* [11] explained the

role of telomeres and telomerase in maintaining chromosome integrity, which is crucial for understanding mutation patterns and genomic stability.

Advanced data structures such as suffix trees, tries, and hash-based indexing [12] have also been widely applied in genomic search systems. These structures provide fast lookup capabilities but often require high memory, making suffix arrays and FM-indexes more practical alternatives for large datasets.

Recent research also explores hybrid and adaptive indexing techniques. Learned index structures [13] combine traditional data structures with machine learning models to improve prediction-based searching. Additionally, compressed data structures such as wavelet trees and Burrows–Wheeler Transform (BWT) [14] are widely used to reduce storage requirements while maintaining efficient query performance.

In addition to exact string matching techniques, heuristic-based algorithms such as BLAST [15] have been widely used for DNA sequence alignment. BLAST employs a seed-and-extend approach, where short exact matches are first identified and then extended to find high-scoring alignments. This method significantly reduces computation time compared to exhaustive dynamic programming techniques, making it highly suitable for large genomic databases. However, BLAST prioritizes speed over exactness and may miss certain alignments, which highlights the need for more accurate and adaptive indexing-based approaches.

Overall, existing literature highlights that while many indexing and search techniques exist, there is still a need for hybrid adaptive systems that balance speed, memory usage, and accuracy. This gap motivates the development of a Hybrid Adaptive DNA Indexing Engine that integrates multiple data structures and algorithms for optimized genomic analysis.

III. SYSTEM ARCHITECTURE

The proposed system follows a hybrid architecture that integrates classical suffix-based indexing with learned prediction models for efficient DNA sequence search. The architecture consists of the following components:

- **Input Module:** Accepts DNA sequence datasets and query patterns.
- **Preprocessing Module:** Cleans and formats the DNA sequence for indexing.
- **Indexing Engine:** Constructs suffix arrays and auxiliary structures such as LCP arrays.
- **Learning Module:** Trains predictive models (e.g., piecewise linear models) to estimate suffix array positions.
- **Search Module:** Performs optimized pattern matching using predicted ranges.

- **Output Module:** Displays matched positions and analysis results.

The integration of prediction models reduces the search space and improves overall efficiency compared to traditional indexing methods.

Hybrid Adaptive DNA Indexing Engine Architecture




Fig 1.1 Hybrid Architecture for DNA sequence Search

III. METHODOLOGY

The proposed methodology integrates classical suffix array–based indexing with predictive modeling techniques to improve the efficiency of DNA sequence search. The system is designed

to reduce search space and computational overhead while maintaining high accuracy. The overall workflow consists of data preprocessing, index construction, model training, and optimized search operations.

A. Data Collection and Preprocessing

DNA sequences are obtained from publicly available genomic datasets and used as input for indexing and evaluation. Since raw genomic data may contain ambiguous or inconsistent characters, a preprocessing step is performed to ensure data quality and uniformity. All sequences are normalized to uppercase format, and invalid or ambiguous nucleotide symbols (such as N, R, Y) are either removed or handled appropriately. To improve memory efficiency and computational performance, DNA sequences are encoded using a compact representation scheme. A common approach is 2-bit encoding, where nucleotides are mapped as follows: A = 00, C = 01, G = 10, and T = 11. This encoding reduces storage requirements and enhances cache locality during processing.

B. Suffix Array Construction

A suffix array is constructed for the given DNA sequence to enable efficient substring search operations. The suffix array is a sorted array of all suffixes of the input sequence arranged in lexicographical order. In this work, an efficient construction algorithm such as the prefix-doubling method or SA-IS algorithm is employed to achieve optimal time complexity. Additionally, the Longest Common Prefix (LCP) array is generated using Kasai's algorithm to store the length of the longest common prefixes between adjacent suffixes. Auxiliary data structures, such as rank arrays or occurrence tables, may also be maintained to support faster query processing.

C. Predictive Model Training

To enhance search efficiency, a predictive model is trained to approximate the mapping between DNA substrings and their corresponding positions in the suffix array. The model treats indexing as a regression problem, where the goal is to estimate the suffix array index for a given query pattern. Training data is generated by extracting prefixes (or k-mers) from the DNA sequence and associating them with their positions in the suffix array. A piecewise linear model or similar regression-based approach is employed to learn this mapping. The model minimizes prediction error using standard loss functions such as mean squared error (MSE). Formally, the model learns a function:

$$f(p) \approx SA_index$$

where p represents the query pattern and SA_index denotes its approximate position in the suffix array.

D. Error Estimation and Correction

Since the predictive model may not provide exact positions, an error bound is estimated to ensure correctness. For each prediction, a maximum error value ϵ is computed based on training data. This allows the system to define a bounded search interval around the predicted position:

$$[f(p) - \epsilon, f(p) + \epsilon]$$

This bounded range guarantees that all possible matches are included, thereby preserving the accuracy of the search process.

E. Search Optimization

Instead of performing a full binary search over the entire suffix array, the predicted position and error bounds are used to significantly reduce the search space. A localized binary search is conducted within the bounded interval to locate candidate matches efficiently.

After identifying candidate positions, a verification step is performed through character-by-character comparison to ensure exact pattern matching. This hybrid approach reduces the time complexity of search operations while maintaining correctness.

F. Algorithm Steps

Algorithm 1: Hybrid DNA Sequence Search

Input: DNA sequence S , query pattern P

Output: Set of matching positions

1. Preprocess the DNA sequence S
2. Construct the suffix array SA and LCP array
3. Train predictive model M using prefix-to-index mapping
4. Predict approximate position: $pos = M(P)$
5. Define search interval: $[pos - \epsilon, pos + \epsilon]$
6. Perform binary search within the interval
7. Verify matches using exact comparison
8. Return all matching positions

G. Experimental Setup and Evaluation

The proposed system is implemented using a combination of C++ (for efficient suffix array construction and search operations) and Python (for model training and visualization). Experiments are conducted on standard genomic datasets of varying sizes to evaluate performance.

The system is evaluated using metrics such as search time, memory usage, and prediction accuracy. Comparative analysis is performed against traditional methods such as standard suffix arrays and FM-indexes to demonstrate improvements in efficiency and scalability.

V. RESULTS AND DISCUSSION

The proposed system was evaluated based on performance metrics such as search time, memory usage, and accuracy.

A. Performance Analysis

The hybrid approach significantly reduces search time compared to traditional suffix array methods. By narrowing the search range, the number of comparisons is minimized.

B. Memory Efficiency

The system maintains a low memory overhead as the predictive model requires minimal additional storage compared to the suffix array.

C. Scalability

The approach performs efficiently on large DNA datasets, demonstrating its ability to handle real-world genomic data.

D. Discussion

The results indicate that integrating learned models with classical indexing techniques improves efficiency without compromising accuracy. However, performance may vary depending on dataset characteristics and model quality.

Base Frequency Distribution




Fig 1.2 Base Frequency Distribution

Nucleotide Composition




Fig 1.3 Nucleotide Composition

GC Content Analysis




Fig 1.4 GC Analysis

VI. CONCLUSION

This paper presents a hybrid DNA indexing and search framework that integrates classical suffix array techniques with learned prediction models to improve the efficiency of substring search in large genomic datasets. By leveraging the strengths of structured indexing and predictive modeling, the proposed approach effectively reduces the search space and minimizes unnecessary memory accesses, which are common bottlenecks

in traditional methods [2], [4]. The use of a predictive model to estimate approximate suffix array positions enables faster query processing while maintaining high accuracy through bounded search and verification mechanisms, similar to learned index-based approaches proposed in recent studies [1], [13]. The experimental evaluation demonstrates that the proposed system achieves significant improvements in query time and overall performance compared to conventional indexing approaches such as standard suffix arrays and FM-indexes [12], [14]. Additionally, the method exhibits better cache utilization and scalability, making it suitable for handling large-scale genomic data generated by modern sequencing technologies [6].

Furthermore, the hybrid architecture provides a flexible framework that can adapt to different types of genomic datasets and query patterns. The integration of machine learning into indexing not only enhances performance but also opens new possibilities for intelligent data structures in bioinformatics [13]. Future work will focus on incorporating more advanced machine learning techniques, such as deep learning and neural network-based models, to further improve prediction accuracy and reduce error bounds. Additionally, efforts will be made to optimize the system for real-time genomic applications, including streaming data processing and large-scale distributed environments. Extending the framework to support approximate matching and mutation-aware searches, which are essential in biological sequence comparison [9], also represents a promising direction for further research.

In conclusion, the proposed hybrid system demonstrates that combining classical algorithms with modern learning-based approaches can significantly advance the efficiency and scalability of DNA sequence search. This work contributes to addressing the growing computational challenges in genomic data analysis and highlights the potential of hybrid intelligent indexing techniques in next-generation bioinformatics systems [1], [15].

VII. REFERENCES

- [1] M. Kirsche, A. Das, and M. C. Schatz, "Sapling: Accelerating suffix array queries with learned data models," *Bioinformatics*, vol. 37, no. 6, pp. 744–749, 2021.
- [2] T. D. Wu, "Bitpacking techniques for indexing genomes: Enhanced suffix arrays," *Algorithms for Molecular Biology*, vol. 11, no. 9, 2016.
- [3] M. Barsky, J. Gabor, M. P. Consens, and A. Thomo, "Suffix Rank: A scalable algorithm for indexing large string collections," *Proceedings of the VLDB Endowment*, vol. 13, no. 11, pp. 2787–2800, 2020.
- [4] Z. Rabea, S. El-Metwally, S. Elmougy, and M. Zakaria, "A fast algorithm for constructing suffix arrays for DNA alphabets," *Journal of King Saud University – Computer and Information Sciences*, vol. 34, pp. 4659–4668, 2022.
- [5] Z.-D. Yang, H.-Y. Kuo, P.-W. Hsieh, and J.-H. Hung, "Efficient construction and utilization of k-ordered FM-indexes," *Bioinformatics*, vol. 40, no. 7, 2024.
- [6] S. M. Wohoush and M. H. Saheb, "Indexing for large DNA database sequences," *International Journal of Biometrics and Bioinformatics*, vol. 5, no. 4, pp. 202–208, 2011.
- [7] D. E. Knuth, J. H. Morris, and V. R. Pratt, "Fast pattern matching in strings," *SIAM Journal on Computing*, vol. 6, no. 2, pp. 323–350, 1977.
- [8] R. S. Boyer and J. S. Moore, "A fast string searching algorithm," *Communications of the ACM*, vol. 20, no. 10, pp. 762–772, 1977.
- [9] V. I. Levenshtein, "Binary codes capable of correcting deletions, insertions, and reversals," *Soviet Physics Doklady*, vol. 10, no. 8, pp. 707–710, 1966.
- [10] Y. Benita, R. S. Oosting, M. C. Lok, M. J. Wise, and I. Humphery-Smith, "Regionalized GC content of template DNA as a predictor of PCR success," *Nucleic Acids Research*, vol. 31, no. 16, 2003.
- [11] E. H. Blackburn, J. W. Szostak, and C. W. Greider, "Maintenance of chromosomes by telomeres and the enzyme telomerase," Nobel Prize Lecture, 2009.
- [12] E. Ukkonen, "On-line construction of suffix trees," *Algorithmica*, vol. 14, no. 3, pp. 249–260, 1995.
- [13] T. Kraska, A. Beutel, E. H. Chi, J. Dean, and N. Polyzotis, "The case for learned index structures," *Proceedings of the ACM SIGMOD*, pp. 489–504, 2018.
- [14] P. Ferragina and G. Manzini, "Opportunistic data structures with applications," *Proceedings of the IEEE FOCS*, pp. 390–398, 2000.
- [15] S. F. Altschul et al., "Basic local alignment search tool," *Journal of Molecular Biology*, vol. 215, no. 3, pp. 403–410, 1990.